

# Data Mover Go application NFR - Non Functional Requirements Q&A.

WA Brown aes@3pp.com

## Is Data-Mover multi-threaded?

Yes, **Data-Mover** is multi-threaded. Both the Pub/Sub subscriber and Kafka producer inherently handle concurrent operations:

### 1. Pub/Sub Subscriber:

- The Subscription.Receive method spawns goroutines to invoke the message handler concurrently for each message.
- The concurrency is managed by the MaxOutstandingMessages parameter in ReceiveSettings.

### 2. Kafka Producer:

- The Kafka producer uses an internal thread pool for sending messages asynchronously.
- Delivery reports are processed in a separate goroutine created within the NewProducer function.

## Will it handle concurrency? How? What number of concurrent calls?

Yes, **Data-Mover** handles concurrency effectively:

### 1. Pub/Sub:

- The Subscription.Receive function supports concurrent message processing by default.
- The concurrency level is determined by the MaxOutstandingMessages setting. In this code, it is set to 100, meaning up to 100 messages can be processed concurrently.

### 2. Kafka:

- Kafka producer supports high concurrency internally by allowing multiple threads to publish messages asynchronously. No explicit concurrency limits are defined in the producer.

### Number of Concurrent Calls:

- **Pub/Sub:** Up to 100 concurrent messages (as configured).
- **Kafka:** Concurrency depends on Kafka's internal configuration and broker capacity, which can handle thousands of messages concurrently if configured properly.

## Will it scale? How? How many jobs per second?

Yes, **Data-Mover** will scale effectively due to its architecture:

### 1. Scaling Mechanisms:

- **Pub/Sub:** Google Pub/Sub can scale dynamically to handle a large

number of messages by increasing the number of subscribers or their MaxOutstandingMessages limits.

- **Kafka:** Kafka is inherently scalable by adding partitions to topics and increasing the number of brokers in the cluster.

## 2. **Horizontal Scaling:**

- The application can scale horizontally by deploying multiple instances of **Data-Mover**, each with its own subscriber and producer.

## 3. **Jobs Per Second:**

- Pub/Sub supports millions of messages per second, so the limit is determined by:
  - ♦ The number of subscribers (MaxOutstandingMessages).
  - ♦ The speed of processing and network throughput.
  - ♦ Kafka broker capacity.

- 4. **Estimate:** With current settings (100 concurrent messages and assuming minimal processing delay), **Data-Mover** could handle hundreds of messages per second per instance, assuming no bottlenecks in Kafka or Pub/Sub.

# Is the current error handling sufficient? If not, what should be added?

## Current Error Handling:

### 1. **Pub/Sub Subscriber:**

- Acks or Nacks messages based on the success of the MessageHandler.
- Logs errors when messages fail to process.

### 2. **Kafka Producer:**

- Logs errors for message delivery failures using Kafka's delivery report mechanism.

### 3. **Bridge:**

- Updates failure metrics for both Pub/Sub and Kafka errors.
- Logs errors for failed message processing.

## Improvements:

### 1. **Retry Mechanism:**

- Add a retry policy for failed messages (both Pub/Sub and Kafka).
- Pub/Sub messages are Nacked, which allows retries by Pub/Sub, but Kafka message failures should include a retry mechanism before logging as a failure.

### 2. **Dead Letter Queue (DLQ):**

- Implement a DLQ for messages that cannot be processed after multiple retries.
- Pub/Sub supports DLQs natively, while Kafka DLQs can be

implemented as a separate topic.

3. **Circuit Breaker Pattern:**

- Prevent cascading failures during high error rates by pausing processing temporarily.

4. **Backpressure Handling:**

- Monitor Kafka producer queue size and slow down Pub/Sub message intake if the Kafka producer cannot keep up.

5. **Metrics Alerts:**

- Add alerting mechanisms for key metrics (e.g., high failure rates, increased processing time).

## **Summary:**

1. **Concurrency:** Data-Mover supports concurrent message processing with both Pub/Sub and Kafka.
2. **Scalability:** It can scale horizontally and handle potentially thousands of jobs per second, limited by resource configuration.
3. **Error Handling:** Current handling is basic but functional. Adding retries, DLQs, circuit breakers, and alerting would make it robust.